

In [6]:

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
pd.set_option('display.max_column',None)
from random import sample
from sklearn import preprocessing
from sklearn.preprocessing import scale
import seaborn as sns
from scipy import stats
import scipy
import statsmodels.api as sm
from statsmodels.formula.api import ols
from statsmodels.stats.multicomp import pairwise_tukeyhsd
import statsmodels.formula.api as smf
from statsmodels.stats.outliers_influence import variance_inflation_factor
pd.set_option('display.max_column',None)
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
import plotly.graph_objects as go
import plotly.io as pio

pio.templates.default = 'plotly_white'
```

In [2]:

```
df=pd.read_csv('Asteroid.csv')
```

```
C:\Users\saina\anaconda3\lib\site-packages\IPython\core\interactiveshell
1.py:3146: DtypeWarning: Columns (11,14,15,22,23) have mixed types.Speci
fy dtype option on import or set low_memory=False.
    has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```

In [3]:

```
df
```

Out[3]:

	full_name	a	e	G	i	om	w	q
0	1 Ceres	2.769165	0.076009	0.12	10.594067	80.305532	73.597694	2.558684
1	2 Pallas	2.772466	0.230337	0.11	34.836234	173.080063	310.048857	2.133865
2	3 Juno	2.669150	0.256942	0.32	12.988919	169.852760	248.138626	1.983332
3	4 Vesta	2.361418	0.088721	0.32	7.141771	103.810804	150.728541	2.151909
4	5 Astraea	2.574249	0.191095	NaN	5.366988	141.576604	358.687608	2.082324
...
839731	(6344 P-L)	2.812945	0.664688	NaN	4.695700	183.310012	234.618352	0.943214
839732	(1168 T-2)	2.645238	0.259376	NaN	12.574937	1.620020	339.568072	1.959126
839733	(2060 T-2)	2.373137	0.202053	NaN	0.732484	176.499082	198.026527	1.893638
839734	(2678 T-3)	2.260404	0.258348	NaN	9.661947	204.512448	148.496988	1.676433
839735	(4571 T-3)	2.546442	0.287672	NaN	5.356238	70.709555	273.483265	1.813901

839736 rows × 27 columns

In [7]:

```
!pip install tensorflow_decision_forests --upgrade -q
```

ERROR: Could not find a version that satisfies the requirement tensorflow_decision_forests (from versions: none)
ERROR: No matching distribution found for tensorflow_decision_forests

In [8]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839736 entries, 0 to 839735
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   full_name             839736 non-null object
1   a                     839734 non-null float64
2   e                     839736 non-null float64
3   G                     119 non-null   float64
4   i                     839736 non-null float64
5   om                    839736 non-null float64
6   w                     839736 non-null float64
7   q                     839736 non-null float64
8   ad                    839730 non-null float64
9   per_y                 839735 non-null float64
10  data_arc              823947 non-null float64
11  condition_code        838743 non-null object
12  n_obs_used            839736 non-null int64
13  H                     837042 non-null float64
14  diameter              137681 non-null object
15  extent                18 non-null    object
16  albedo                136452 non-null float64
17  rot_per               18796 non-null float64
18  GM                    14 non-null    float64
19  BV                    1021 non-null  float64
20  UB                    979 non-null   float64
21  IR                     1 non-null     float64
22  spec_B                1666 non-null  object
23  spec_T                980 non-null   object
24  neo                   839730 non-null object
25  pha                   822814 non-null object
26  moid                  822814 non-null float64
dtypes: float64(18), int64(1), object(8)
memory usage: 173.0+ MB
```

In [9]:

```
df = df.dropna(subset=['diameter'])
```

In [10]:

df.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 137681 entries, 0 to 810411
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   full_name             137681 non-null object
1   a                     137681 non-null float64
2   e                     137681 non-null float64
3   G                     119 non-null    float64
4   i                     137681 non-null float64
5   om                    137681 non-null float64
6   w                     137681 non-null float64
7   q                     137681 non-null float64
8   ad                    137681 non-null float64
9   per_y                 137681 non-null float64
10  data_arc               137541 non-null float64
11  condition_code         137681 non-null object
12  n_obs_used              137681 non-null int64
13  H                      136930 non-null float64
14  diameter                137681 non-null object
15  extent                  16 non-null     object
16  albedo                  136449 non-null float64
17  rot_per                 11188 non-null  float64
18  GM                      14 non-null     float64
19  BV                      1005 non-null   float64
20  UB                      965 non-null    float64
21  IR                       1 non-null      float64
22  spec_B                  1370 non-null   object
23  spec_T                   965 non-null    object
24  neo                     137681 non-null object
25  pha                     137681 non-null object
26  moid                    137681 non-null float64
dtypes: float64(18), int64(1), object(8)
memory usage: 29.4+ MB

```

In [11]:

```
df = df.drop(columns=['full_name', 'H', 'albedo', 'G', 'extent', 'rot_per', 'GM', 'BV'],
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 137681 entries, 0 to 810411
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      137681 non-null float64
1   e                      137681 non-null float64
2   i                      137681 non-null float64
3   om                     137681 non-null float64
4   w                      137681 non-null float64
5   q                      137681 non-null float64
6   ad                     137681 non-null float64
7   per_y                 137681 non-null float64
8   data_arc              137541 non-null float64
9   condition_code        137681 non-null object
10  n_obs_used             137681 non-null int64
11  diameter               137681 non-null object
12  neo                    137681 non-null object
13  pha                    137681 non-null object
14  moid                   137681 non-null float64
dtypes: float64(10), int64(1), object(4)
memory usage: 16.8+ MB
```

In [12]:

```
df.diameter.dtype
```

Out[12]:

```
dtype('O')
```

In [13]:

```
df.diameter = pd.to_numeric(df.diameter, errors='coerce')
```

In [14]:

```
np.sum(df.diameter.isna())
```

Out[14]:

```
1
```

In [15]:

```
df = df.dropna(subset=['diameter'])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 137680 entries, 0 to 810411
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      137680 non-null float64
1   e                      137680 non-null float64
2   i                      137680 non-null float64
3   om                     137680 non-null float64
4   w                      137680 non-null float64
5   q                      137680 non-null float64
6   ad                     137680 non-null float64
7   per_y                 137680 non-null float64
8   data_arc              137540 non-null float64
9   condition_code        137680 non-null object
10  n_obs_used             137680 non-null int64
11  diameter               137680 non-null float64
12  neo                    137680 non-null object
13  pha                    137680 non-null object
14  moid                   137680 non-null float64
dtypes: float64(11), int64(1), object(3)
memory usage: 16.8+ MB
```

In [16]:

```
df.diameter.describe()
```

Out[16]:

```
count    137680.000000
mean         5.480873
std         9.365499
min         0.002500
25%         2.770000
50%         3.956000
75%         5.741000
max        939.400000
Name: diameter, dtype: float64
```

In [17]:

```
df = df.dropna(subset=['data_arc'])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 137540 entries, 0 to 810411
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      137540 non-null float64
1   e                      137540 non-null float64
2   i                      137540 non-null float64
3   om                    137540 non-null float64
4   w                    137540 non-null float64
5   q                    137540 non-null float64
6   ad                   137540 non-null float64
7   per_y               137540 non-null float64
8   data_arc            137540 non-null float64
9   condition_code      137540 non-null object
10  n_obs_used          137540 non-null int64
11  diameter            137540 non-null float64
12  neo                 137540 non-null object
13  pha                 137540 non-null object
14  moid                137540 non-null float64
dtypes: float64(11), int64(1), object(3)
memory usage: 16.8+ MB
```

In [18]:

```
df.condition_code.value_counts()
```

Out[18]:

```
0      126192
9       5584
0       2017
9.0      827
1       809
1       563
2       302
5       242
6       152
4       144
3       135
7       129
2        95
5.0      92
7.0      78
8        49
6.0      36
4.0      33
3        33
8.0      28
Name: condition_code, dtype: int64
```

In [19]:

```
df.condition_code = pd.to_numeric(df.condition_code, errors='coerce')
```

In [20]:

```
df.neo.value_counts()
```

Out[20]:

```
N    136691
Y      849
Name: neo, dtype: int64
```

In [21]:

```
df.neo = df.neo.replace({'N':0, 'Y':1})
```

In [22]:

```
df.pha.value_counts()
```

Out[22]:

```
N    137320
Y      220
Name: pha, dtype: int64
```

In [23]:

```
df.pha = df.pha.replace({'N':0, 'Y':1})
```


In [24]:

```
df.info()
```

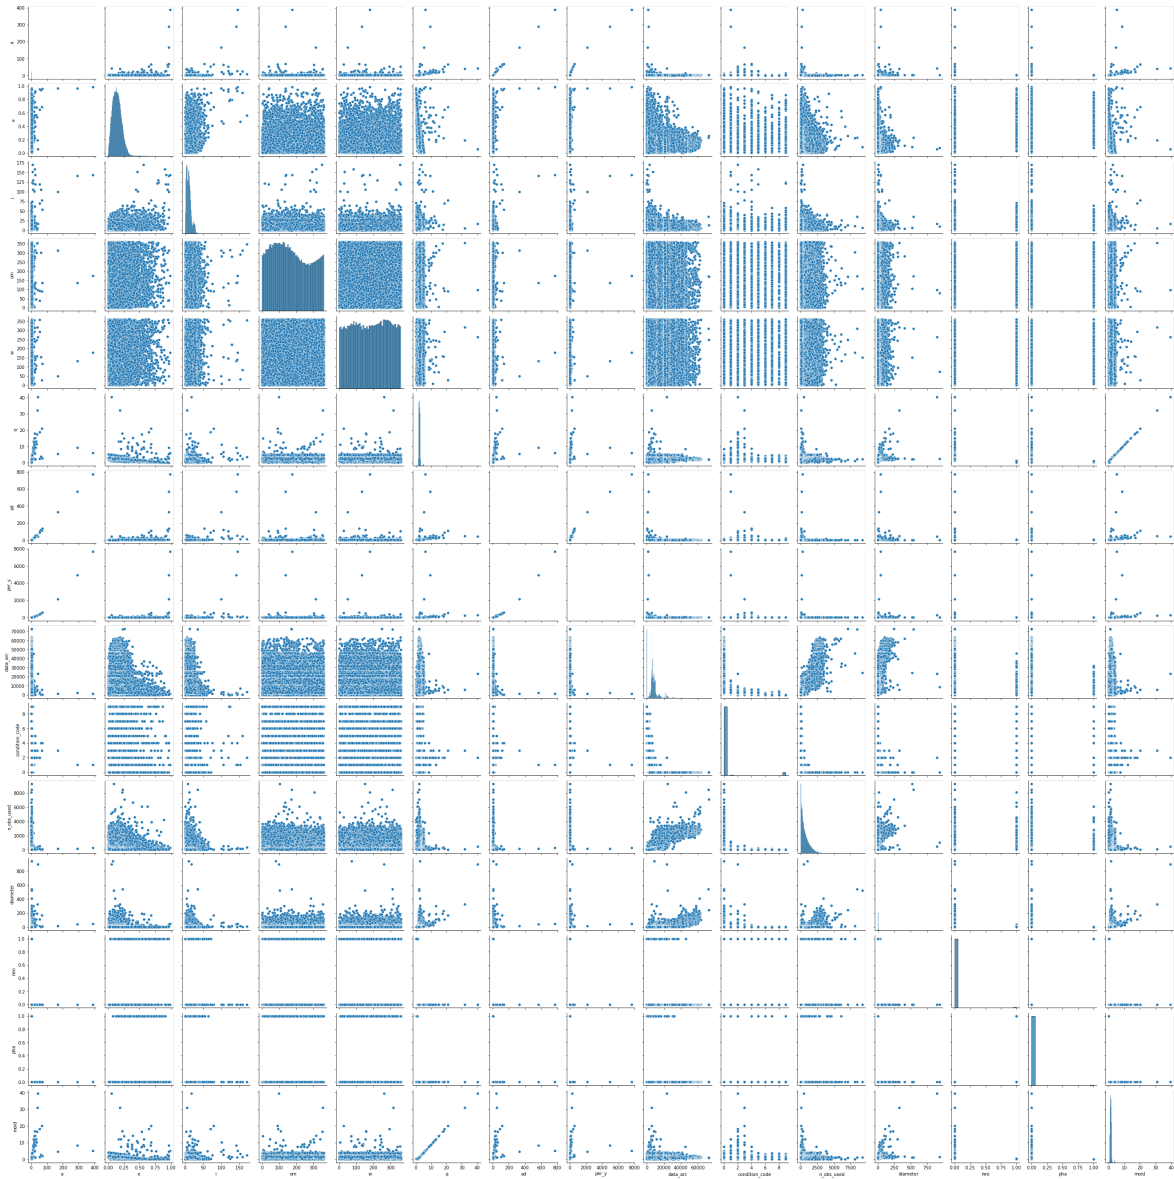
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 137540 entries, 0 to 810411
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   a                      137540 non-null float64
1   e                      137540 non-null float64
2   i                      137540 non-null float64
3   om                     137540 non-null float64
4   w                      137540 non-null float64
5   q                      137540 non-null float64
6   ad                     137540 non-null float64
7   per_y                 137540 non-null float64
8   data_arc              137540 non-null float64
9   condition_code        137540 non-null float64
10  n_obs_used             137540 non-null int64
11  diameter               137540 non-null float64
12  neo                    137540 non-null int64
13  pha                    137540 non-null int64
14  moid                   137540 non-null float64
dtypes: float64(12), int64(3)
memory usage: 16.8 MB
```

In [37]:

```
sns.pairplot(df)
```


Out[37]:

<seaborn.axisgrid.PairGrid at 0x17c85a7f760>



In [57]:

```
plt.figure(figsize=(25,12),dpi=150)
sns.heatmap(df.corr(),annot=True,cmap='YlGnBu',lw=0.2,annot_kws={'size':20})
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)
plt.show()
```



In [25]:

```
features = df.drop('diameter', axis=1)
target = df.diameter
```

In [26]:

```
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, ra
```

In [27]:

```
forest = RandomForestRegressor(max_depth=32,n_estimators=50)
```

In [28]:

```
for column in df.columns:  
    print(column, np.sum(df[column].isna()))
```

```
a 0  
e 0  
i 0  
om 0  
w 0  
q 0  
ad 0  
per_y 0  
data_arc 0  
condition_code 0  
n_obs_used 0  
diameter 0  
neo 0  
pha 0  
moid 0
```

In []:

```
#Random Forsert Regressor
```

In [34]:

```
forest.fit(X_train, np.ravel(y_train))
```

Out[34]:

```
RandomForestRegressor(max_depth=32, n_estimators=50)
```

In [35]:

```
y_pred = forest.predict(X_test)
```

In [36]:

```
r2_score(y_test, y_pred)
```

Out[36]:

```
0.804780635717621
```

In [55]:

```
from sklearn import metrics  
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))  
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))  
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
Mean Absolute Error: 1.2124214970776614  
Mean Squared Error: 15.0206748969232  
Root Mean Squared Error: 3.8756515448274245
```

In [56]:

```
df.diameter.mean()
```

Out[56]:

5.483639584848152

In []:

```
# it looks like my model is good because rmse is less than target variable mean
```